

# Boosting algorithms for Supervised Learning

## A practical study of *AdaBoost*

Alexandre Combessie    Ismail Machraoui

ENSAE ParisTech

Project for the "Aggregation in Statistical Learning" course  
by Pierre Alquier, 2016

- Freund and Schapire (1996) start with a generalized algorithm for the **on-line allocation model**
  - Example: allocating bets among horse-racing "experts"
  - Adaptation of Littlestone and Warmuth (1994) multiplicative weight-update rule for majority voting
- From this on-line setting, they derive a **Boosting algorithm** for **supervised, batch learning**
  - Idea: convert a family of "Weak Learner" algorithms into a single "Strong Learner"
  - Non-obvious reversal of the online-to-batch framework
- Starting with a simple algorithm for **binary classification**, they extend it to **multi-class classification** and **regression**

## Boosting algorithms for Supervised Learning

Alexandre  
Combesse,  
Ismail  
Machraoui

## Introduction

### Algorithm description

On-line  
allocation origin  
Batch binary  
classification  
Multi-class  
extensions

### Mathematical results

Training  
Generalization

### Empirical study

Chosen  
implementations  
Performance  
analysis

- 1 Algorithm description
  - On-line allocation origin
  - Batch binary classification
  - Multi-class extensions
- 2 Mathematical results
  - Training
  - Generalization
- 3 Empirical study
  - Chosen implementations
  - Performance analysis

- Starting example: How to allocate money among gamblers on horse races? (a bit like bandits, but with agents)
- Mathematical framework:
  - Agent  $A$  with  $\{1, \dots, N\}$  *strategies* to choose from
  - At each *trial*  $t = \{1, \dots, T\}$ , agent  $A$  decides on the distribution  $\mathbf{p}^t$  over the strategies
  - Each strategy  $i$  incurs loss  $l_i^t$  (possibly in adversarial environment) so that the loss of  $A$  is  $\mathbf{p}^t \cdot \mathbf{l}^t$
  - The goal of agent  $A$  is to minimize its cumulative loss compared to the loss of the best strategy:

$$\min_{\mathbf{p}^t} \left\{ \sum_{t=1}^T \mathbf{p}^t \cdot \mathbf{l}^t - \min_i \sum_{t=1}^T l_i^t \right\}$$

# Proposed on-line allocation algorithm

## Hedge( $\beta$ )

Boosting  
algorithms for  
Supervised  
Learning

Alexandre  
Combesse,  
Ismail  
Machraoui

Introduction

Algorithm  
description

On-line  
allocation origin

Batch binary  
classification

Multi-class  
extensions

Mathematical  
results

Training  
Generalization

Empirical  
study

Chosen  
implementations  
Performance  
analysis

### Initialization:

Weight multiple parameter  $\beta \in [0, 1]$

Number of trials  $T \in \mathbb{N}^*$

Initial weight vector  $\mathbf{w}^1 \in [0, 1]^N$  with  $\sum_{i=1}^N w_i^1 = 1$

**for**  $t = 1$  to  $T$  number of trial **do**

Set strategy allocation  $\mathbf{p}^t \leftarrow \frac{\mathbf{w}^t}{\sum_{i=1}^N w_i^t}$

Observe realization with loss vector  $\mathbf{l}^t$

Incur loss of  $\mathbf{p}^t \cdot \mathbf{l}^t$

**for**  $i = 1$  to  $N$  **do**

Update weight vector  $w_i^{t+1} \leftarrow w_i^t \times \beta^{l_i^t}$

**end for**

**end for**

# Preliminary framework to define the concepts of Boosting, Weak Learners and Strong Learners

## Boosting algorithms for Supervised Learning

Alexandre  
Combesse,  
Ismail  
Machraoui

## Introduction

## Algorithm description

On-line  
allocation origin

Batch binary  
classification

Multi-class  
extensions

## Mathematical results

Training  
Generalization

## Empirical study

Chosen  
implementations  
Performance  
analysis

- Boosting is the transformation of a "Weak Learner" algorithm into a "Strong Learner" one
- The notion of "Weak Learner" is defined in the Probably Approximately Correct (PAC) learning framework
  - A "Strong Learner" is an algorithm that given  $\varepsilon, \delta > 0$  outputs a hypothesis with error  $< \varepsilon$  with probability  $1 - \delta$
  - A "Weak Learner" only verifies it for  $\varepsilon \geq 1/2 - \gamma$  where  $\gamma > 0$  or decreases as  $1/p$  with  $p$  polynomial
  - For simplification, the PAC learning framework is not used in the rest of the paper, in favor of a more general one where examples  $(x_i, y_i)$  are chosen randomly according to a fixed but unknown distribution  $\mathfrak{P}$  over  $X \times Y$
- In the context of *batch* learning we will focus on boosting by *sampling* over the examples

# How to adapt this on-line allocation algorithm to boosting problems in batch settings? (1/2)

## Boosting algorithms for Supervised Learning

Alexandre  
Combesse,  
Ismail  
Machraoui

## Introduction

## Algorithm description

On-line  
allocation origin

**Batch binary  
classification**

Multi-class  
extensions

## Mathematical results

Training  
Generalization

## Empirical study

Chosen  
implementations  
Performance  
analysis

- The original boosting algorithm was developed by Schapire, and improved by Freund as the "boost-by-majority" algorithm
- The issue of these algorithms is that they require to know the bias of the Weak Learner in advance, and does not make us of all the Weak Learner hypothesis.
- **To solve these problems, Freund and Schapire decide to adapt their online allocation algorithm in the context of boosting**

# How to adapt this on-line allocation algorithm to boosting problems in batch settings? (2/2)

## Boosting algorithms for Supervised Learning

Alexandre  
Combesse,  
Ismail  
Machraoui

### Introduction

### Algorithm description

### On-line allocation origin

### Batch binary classification

### Multi-class extensions

### Mathematical results

### Training Generalization

### Empirical study

### Chosen implementations Performance analysis

- There are correspondences between the on-line allocation model and the problem of boosting
- The authors actually choose the less obvious reverse correspondence:

On-line allocation	Boosting problem	
	<i>Natural</i>	<i>Reversed</i>
Strategy	Weak Learner	Training example
Trial	Training example	Weak Learner



# Boosting algorithm for binary $Y = \{0, 1\}$

## AdaBoost

Boosting  
algorithms for  
Supervised  
Learning

Alexandre  
Combesse,  
Ismail  
Machraoui

Introduction

Algorithm  
description

On-line  
allocation origin

Batch binary  
classification

Multi-class  
extensions

Mathematical  
results

Training  
Generalization

Empirical  
study

Chosen  
implementations  
Performance  
analysis

**Input:** Labeled examples  $(x_1, y_1) \dots (x_N, y_N)$  with distribution  $D$ , algorithm **WeakLearn**, number of iterations  $T$

**Initialization:** Weight vector  $w_i^1 \leftarrow D(i)$  **for**  $i = 1$  to  $N$   
**for**  $t = 1$  to  $T$  number of iterations **do**

Set example distribution  $\mathbf{p}^t \leftarrow \frac{\mathbf{w}^t}{\sum_{i=1}^N w_i^t}$

Call **WeakLearn** with distribution  $\mathbf{p}^t$ , get hypothesis  $h_t$

Compute  $h_t$  loss  $\varepsilon_t \leftarrow \sum_{i=1}^N p_i^t \mathbb{I}[h_t(x_i) \neq y_i]$

Set  $\beta_t \leftarrow \varepsilon_t / (1 - \varepsilon_t)$

Update weight  $w_i^{t+1} \leftarrow w_i^t \times \beta^{1 - \mathbb{I}[h_t(x_i) \neq y_i]}$  **for**  $i = 1$  to  $N$

**end for**

**return** Final hypothesis

$$h_f(x) = \operatorname{argmax}_{y \in Y} \sum_{t=1}^T \left( \log \frac{1}{\beta_t} \right) \mathbb{I}[h_t(x) \neq y]$$

# Multi-class extensions $Y = \{1, \dots, k\}$

Boosting  
algorithms for  
Supervised  
Learning

Alexandre  
Combesse,  
Ismail  
Machraoui

Introduction

Algorithm  
description

On-line  
allocation origin

Batch binary  
classification

Multi-class  
extensions

Mathematical  
results

Training  
Generalization

Empirical  
study

Chosen  
implementations  
Performance  
analysis

There are several ways to extend the previous algorithm in multi-class contexts:

- ① **AdaBoost.M1**: Naive extension of binary AdaBoost by replacing  $Y = \{0, 1\}$  by  $Y = \{1, \dots, k\}$  and adding a rule to abort the main loop if  $\varepsilon_t > 1/2$
- ② **Adaboost.M2**: Advanced extension of AdaBoost with more communication between the boosting and Weak Learner algorithm: probabilities and class weights
- ③ **Binarization**: For  $k$  classes, perform boosting separately on the  $k$  binarized problems and aggregate them according to rules like One-vs-Rest or Error-Correcting Output

- The algorithm is very close to the binary AdaBoost, but it adds an important **loop breaking rule**:

$$\varepsilon_t = \sum_{i=1}^N p_i^t \mathbb{I}[h_t(x_i) \neq y_i] > 1/2$$

- Without this rule, we could have a weight multiplier  $\beta_t > 1$  and the algorithm would diverge
- Inherent problem with this approach: The more  $k$  classes, the harder it is for a Weak Learner to ensure  $\varepsilon_t \leq 1/2$

Modifications to AdaBoost.M1 are highlighted in **red**

**Input:** Labeled examples  $(x_1, y_1) \dots (x_N, y_N)$  with distribution  $D$ ,  
algorithm **WeakLearn**, number of iterations  $T$

**Initialization:**  $w_{i,y}^1 \leftarrow D(i)/(k-1)$  for  $i = 1$  to  $N$  and  $y \in Y - \{y_i\}$   
for  $t = 1$  to  $T$  number of iterations **do**

Set example distribution  $D_t(i) \leftarrow \frac{\sum_{y \neq y_i} w_{i,y}^t}{\sum_{i=1}^N \sum_{y \neq y_i} w_{i,y}^t}$

Set label weights  $q_t(i, y) \leftarrow \frac{w_{i,y}^t}{\sum_{y \neq y_i} w_{i,y}^t}$

Call **WeakLearn** with example distribution  $D_t$  and label weights  $q_t$ ,  
get hypothesis  $h_t$  with probability values in  $[0, 1]$

Compute  $h_t$  pseudo-loss

$\varepsilon_t \leftarrow \frac{1}{2} \sum_{i=1}^N D_t(i) \left( 1 - h_t(x_i, y_i) + \sum_{y \neq y_i} q_t(i, y) h_t(x_i, y) \right)$

Set  $\beta_t \leftarrow \varepsilon_t / (1 - \varepsilon_t)$

$w_{i,y}^{t+1} \leftarrow w_{i,y}^t \times \beta^{\frac{1}{2}(1+h_t(x_i, y_i) - h_t(x_i, y))}$  for  $i = 1$  to  $N$  and  $y \in Y - \{y_i\}$

**end for**

**return** Final hypothesis  $h_f(x) = \operatorname{argmax}_{y \in Y} \sum_{t=1}^T \left( \log \frac{1}{\beta_t} \right) h_t(x, y)$

# Binarization - One vs all

**Input:** Labeled examples  $(x_1, y_1) \dots (x_N, y_N)$  with  $y_i \in \{1, \dots, k\}$  and with distribution  $D$ , algorithm **AdaBoost**, number of iterations  $T$

**for**  $j = 1$  to  $k$  number of classes **do**

Transform  $(x_1, y_1) \dots (x_N, y_N)$  into a binary dataset

$(x_1, y_1^{(j)}) \dots (x_N, y_N^{(j)})$  where  $y_i^{(j)} = \mathbb{I}[y_i = j]$

Call **AdaBoost** on  $(x_1, y_1^{(j)}) \dots (x_N, y_N^{(j)})$  with distribution  $D$  and number of iterations  $T$

Get hypothesis  $h_f^{(j)}$  with probability values in  $[0, 1]$

**end for**

**return** Final hypothesis

$$h_f(x) = \operatorname{argmax}_{j \in \{1, \dots, k\}} h_f^{(j)}(x)$$

# Training error : Adaboost.M1

## Boosting algorithms for Supervised Learning

Alexandre  
Combesse,  
Ismail  
Machraoui

### Introduction

### Algorithm description

On-line  
allocation origin  
Batch binary  
classification  
Multi-class  
extensions

### Mathematical results

Training  
Generalization

### Empirical study

Chosen  
implementations  
Performance  
analysis

- Let  $\varepsilon_1 \dots \varepsilon_t$  be the **Weak Learner** algorithm generated errors :  $\varepsilon_t = \sum_{i=1}^N p_i^t \mathbb{I}[h_t(x_i) \neq y_i]$
- Let  $\varepsilon$  the error of  $h_f$  (the final hypothesis) :

$$\varepsilon \leq 2^T \prod_{t=1}^T \sqrt{\varepsilon_t(1 - \varepsilon_t)} \text{ (similar to **Adaboost**)}$$

Or if  $\varepsilon = \frac{1}{2} - \gamma$  :

$$\varepsilon \leq \prod_{t=1}^T \sqrt{1 - 4\gamma^2}$$

If all the errors of all the hypotheses are equal, the inequality can be simplified to :

$$\varepsilon \leq \exp(-2T\gamma^2)$$

# Training error : Adaboost.M2

## Boosting algorithms for Supervised Learning

Alexandre  
Combesse,  
Ismail  
Machraoui

### Introduction

### Algorithm description

On-line  
allocation origin  
Batch binary  
classification  
Multi-class  
extensions

### Mathematical results

Training  
Generalization

### Empirical study

Chosen  
implementations  
Performance  
analysis

- Let  $\varepsilon_1 \dots \varepsilon_t$  be the **Weak Learner** algorithm generated errors. As a reminder,  $\varepsilon_t$  stands here for the p-loss.
- Let  $\varepsilon$  the error of  $h_f$  (the final hypothesis) :

$$\varepsilon \leq 2^T (k-1) \prod_{t=1}^T \sqrt{\varepsilon_t (1 - \varepsilon_t)}$$

- It can be explained by a reduction to a binary **Adaboost** and then apply **Adaboost** upper-bound to get to the stated result.

# Generalization error

## Boosting algorithms for Supervised Learning

Alexandre  
Combessie,  
Ismail  
Machraoui

### Introduction

### Algorithm description

On-line  
allocation origin  
Batch binary  
classification  
Multi-class  
extensions

### Mathematical results

Training  
Generalization

### Empirical study

Chosen  
implementations  
Performance  
analysis

- Provided that hypotheses of WeakLearn are from a class of functions which VC-dimension  $d \geq 2$  then :  
For any  $\delta > 0$

$$\Pr[|\varepsilon_f - \hat{\varepsilon}| > 2\sqrt{\frac{d_f(\ln(\frac{2N}{d_f}) + \ln(\frac{9}{\delta}))}{N}}] \leq \delta$$

where  $d_f \leq 2(d+1)(T+1)\log_2[e(T+1)]$ ,  $\varepsilon_f$  stands for the generalization error of the final hypothesis

- The cross-validation technique might be used in order to get an optimal performance (i.e smallest error of the final hypothesis  $h_f^T$ ).
- Making the parameter  $T$  vary, the best  $T$  kept is the one minimizing the error of the final hypothesis on the validation set.



# Chosen dataset

## Heart Disease UCI

Boosting  
algorithms for  
Supervised  
Learning

Alexandre  
Combesse,  
Ismail  
Machraoui

Introduction

Algorithm  
description

On-line  
allocation origin

Batch binary  
classification

Multi-class  
extensions

Mathematical  
results

Training  
Generalization

Empirical  
study

Chosen  
implementations

Performance  
analysis

- Heart disease UCI dataset gathers data about 303 patients in Hungary, Switzerland and the USA. It is available on [archive.ics.uci.edu/ml/datasets/Heart+Disease](http://archive.ics.uci.edu/ml/datasets/Heart+Disease).
- The dataset contains 14 features either categorical or numerical, with few missing values.
- The target variable refers to the existence of heart disease for a given patient. It takes values in  $\{0,1,2,3,4\}$  where 0 stands for pathology absence and the other remaining values points out some level degree of the disease.

# Algorithm implementations

Boosting  
algorithms for  
Supervised  
Learning

Alexandre  
Combessie,  
Ismail  
Machraoui

Introduction

Algorithm  
description

On-line  
allocation origin  
Batch binary  
classification  
Multi-class  
extensions

Mathematical  
results

Training  
Generalization

Empirical  
study

Chosen  
implementations  
Performance  
analysis

The algorithms we've chosen to implement are :

- 1 Adaboost.M1
- 2 Adaboost.M2
- 3 Binarization - one vs all

We set max iterations to  $T = 100$

# Performance analysis : Adaboost.M1 (1/2)

## Boosting algorithms for Supervised Learning

Alexandre Combessie,  
Ismail Machraoui

### Introduction

### Algorithm description

On-line allocation origin  
Batch binary classification  
Multi-class extensions


### Mathematical results

Training  
Generalization

### Empirical study

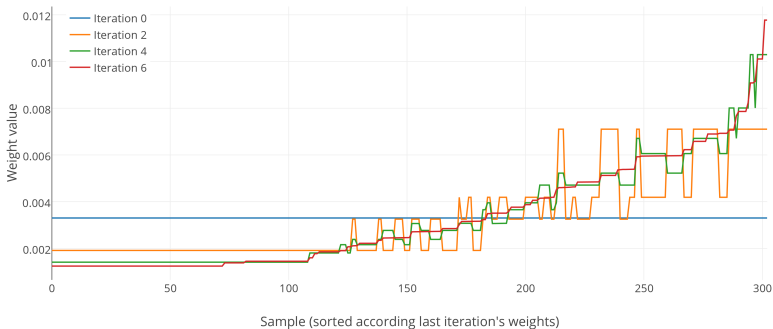
Chosen implementations  
Performance analysis

Evolution of the error rates for AdaBoost.M1



# Performance analysis : Adaboost.M1 (2/2)

Evolution of the sample weight distribution for AdaBoost.M1



# Performance analysis : Adaboost.M2 (1/3)

Boosting  
algorithms for  
Supervised  
Learning

Alexandre  
Combesse,  
Ismail  
Machraoui

Introduction

Algorithm  
description

On-line  
allocation origin

Batch binary  
classification

Multi-class  
extensions

Mathematical  
results

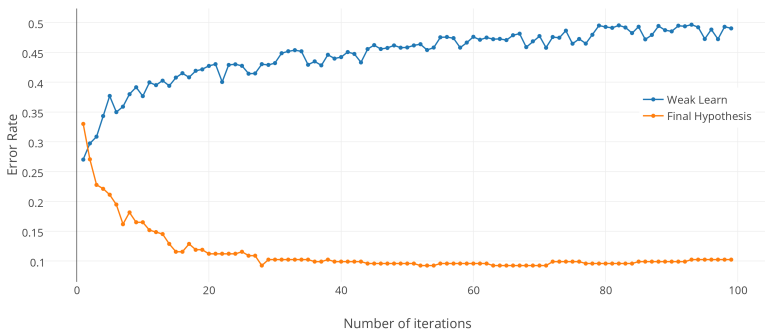
Training  
Generalization

Empirical  
study

Chosen  
implementations


Performance  
analysis

Evolution of the error rates for AdaBoost.M2




# Performance analysis : Adaboost.M2 (2/3)

Evolution of the sample weight distribution for AdaBoost.M2








# Performance analysis : Adaboost.M2 (3/3)

Evolution of the class weight distribution for AdaBoost.M2



# Performance analysis : One vs all

Final Hypothesis error for AdaBoost.M1





# Performance analysis

## Summary

### Boosting algorithms for Supervised Learning

Alexandre Combesse,  
Ismail Machraoui

Introduction

Algorithm description

On-line allocation origin  
Batch binary classification  
Multi-class extensions

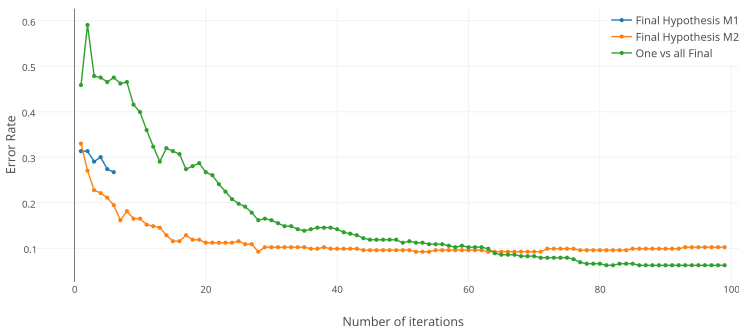
Mathematical results

Training  
Generalization

Empirical study

Chosen implementations  
Performance analysis

### Comparison of the 3 algorithms





Y. Freund and R. Schapire.

A Decision-Theoretic Generalization of On-Line Learning  
and an Application to Boosting

*Journal of Computer and System Sciences*, 55:119–139,  
1997.



Y. Freund and R. Schapire.

A Short Introduction to Boosting

*In Proceedings of the Sixteenth International Joint  
Conference on Artificial Intelligence*, 1401–1406, 1999.